

```
In [6]: import pandas as pd
from matplotlib import pyplot as plt
import seaborn as sns
```

```
In [7]: df = pd.read_csv('insulin_dosage_prediction.csv').drop("patient_id", axis=1)
df
```

Out[7]:

	gender	age	family_history	glucose_level	physical_activity	food_intake	previous_m
0	male	68	yes	103.49	9.28	high	
1	female	57	yes	113.35	6.67	high	
2	male	24	yes	127.40	4.14	medium	
3	male	49	yes	138.79	6.80	medium	
4	male	65	no	128.42	4.01	low	
...	
9995	female	33	no	193.35	6.40	high	
9996	male	77	yes	140.01	2.28	low	
9997	female	71	yes	190.90	8.53	medium	
9998	female	33	yes	164.27	7.50	high	
9999	male	52	yes	104.39	5.17	high	

10000 rows × 14 columns



```
In [33]: print(f"Number of missing values:\n{df.isna().sum()}\nNumber of duplicated rows: {d
```

Number of missing values:

gender	0
age	0
family_history	0
glucose_level	0
physical_activity	0
food_intake	0
previous_medications	0
BMI	0
HbA1c	0
weight	0
insulin_sensitivity	0
sleep_hours	0
creatinine	0
Insulin	0

dtype: int64

Number of duplicated rows: 0

```
In [34]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   gender            10000 non-null   object  
 1   age                10000 non-null   int64  
 2   family_history     10000 non-null   object  
 3   glucose_level      10000 non-null   float64 
 4   physical_activity  10000 non-null   float64 
 5   food_intake        10000 non-null   object  
 6   previous_medications 10000 non-null   object  
 7   BMI                10000 non-null   float64 
 8   HbA1c              10000 non-null   float64 
 9   weight              10000 non-null   float64 
 10  insulin_sensitivity 10000 non-null   float64 
 11  sleep_hours        10000 non-null   float64 
 12  creatinine          10000 non-null   float64 
 13  Insulin             10000 non-null   object  
dtypes: float64(8), int64(1), object(5)
memory usage: 1.1+ MB
```


```
In [35]: # for col in df.columns:
#     print(col)
#     print(np.sort(df[col].unique()))
```

```
In [36]: fig, ax = plt.subplots(2, 7, figsize=(35,10))

colours = sns.color_palette("husl", 14)

for i, col in enumerate(df.columns):
    sns.histplot(df[col], ax=ax[i//7, i%7], color=colours[i])
    ax[i//7, i%7].set_title(col)

plt.suptitle("The Distribution of Various Variables In The Insulin Dosage Dataset")
plt.tight_layout()
plt.show()
```




```
In [74]: fig, ax = plt.subplots(2, 7, figsize=(35,10))


colours = sns.color_palette("husl", 14)

for i, col in enumerate(df.columns):
    sns.boxplot(df[col], ax=ax[i//7, i%7], color=colours[i])
    ax[i//7, i%7].set_title(col)
```

```
plt.suptitle("The Distribution of Various Variables In The Insulin Dosage Dataset")
plt.tight_layout()
plt.show()
```



```
In [37]: fig, ax = plt.subplots(figsize=(12, 12))
sns.heatmap(df.select_dtypes('number').corr(), annot=True)
plt.title("Correlation between various numeric variables of the dataset")
plt.show()
```



```
In [10]: from sklearn.preprocessing import StandardScaler, OneHotEncoder, OrdinalEncoder
from sklearn.compose import ColumnTransformer

from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split, GridSearchCV
```

```
In [11]: model = RandomForestClassifier(n_estimators=200)

num_cols = ['age', 'glucose_level', 'physical_activity', 'BMI', 'HbA1c', 'weight',
nominal_cols = ['gender', 'previous_medications']
ordinal_cols = ['family_history', 'food_intake']

preprocessor = ColumnTransformer(
    transformers=[

        ('num', StandardScaler(), num_cols),
        ('nom', OneHotEncoder(sparse_output=False), nominal_cols),
```

```

        ('ord', OrdinalEncoder(), ordinal_cols)
    ]
)

param_grid = {
    'criterion': ['gini', 'entropy'],
    'n_estimators': [50, 100, 200],
    'max_depth': [None, 10, 20, 30],
    'min_samples_split': [1, 2, 3, 5, 7, 10]
}

```

```
In [12]: x = df.drop('Insulin', axis=1)
y = df['Insulin']

x_transformed = preprocessor.fit_transform(x)

x_train, x_test, y_train, y_test = train_test_split(x_transformed, y, test_size=0.3)
```

```
In [ ]: # grid_search = GridSearchCV(estimator=model, param_grid=param_grid, cv=5, n_jobs=-1)
# grid_search.fit(x_train, y_train)

# just so i don't keep accidentally re-running this tedious bit again and again
```

```
In [42]: from sklearn.metrics import accuracy_score, classification_report, ConfusionMatrixDisplay


best_rf = grid_search.best_estimator_
test_acc = best_rf.score(x_test, y_test)
y_pred = best_rf.predict(x_test)
accuracy_score(y_test, y_pred)
```

Out[42]: 0.978

```
In [43]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
down	0.98	0.87	0.92	190
no	0.98	0.99	0.99	189
steady	0.97	1.00	0.98	1604
up	0.99	0.96	0.98	1017
accuracy			0.98	3000
macro avg	0.98	0.96	0.97	3000
weighted avg	0.98	0.98	0.98	3000

```
In [44]: ConfusionMatrixDisplay.from_predictions(y_test, y_pred)
plt.grid(False)
plt.show()
```



```
In [2]: best_rf = model
```

```
In [40]: from sklearn.utils.validation import check_is_fitted
print(check_is_fitted(best_rf))
```

None

```
In [49]: label_binarizer.classes_
```

```
Out[49]: array(['down', 'no', 'steady', 'up'], dtype='<U6')
```

```
In [81]: # import matplotlib.pyplot as plt
# from sklearn.metrics import RocCurveDisplay
# from sklearn.preprocessing import LabelBinarizer
# y_score = best_rf.predict_proba(x_test)
# label_binarizer = LabelBinarizer().fit(y_train)
# y_onehot_test = label_binarizer.transform(y_test)

# def graph_roc(var):
#     class_id = np.flatnonzero(label_binarizer.classes_ == var)[0]

#     display = RocCurveDisplay.from_predictions(
#         y_onehot_test[:, class_id],
#         y_score[:, class_id],
#         name=f"{var} VS not {var}",
```

```

#         plot_chance_level=True,
#         despine=True,
#     )
#     _ = display.ax_.set(
#         xlabel="False Positive Rate",
#         ylabel="True Positive Rate",
#         title=f"One-vs-Rest ROC curves:{var} vs not {var}",
#     )

# graph_roc('up')
# graph_roc('down')
# graph_roc('steady')
# graph_roc('no')

```

In [45]: `import joblib`

```
joblib.dump(best_rf, 'insulin dosage predictor v2.joblib')
```

Out[45]: `['insulin dosage predictor v2.joblib']`

In []: *## all the stuff below this was for debugging and serves no purpose other than show*

In [1]: `import joblib`

```
model = joblib.load("insulin dosage predictor v2.joblib")
```

In [47]: `transformed_columns = preprocessor.get_feature_names_out()`
`transformed_columns`

Out[47]: `array(['num_age', 'num_glucose_level', 'num_physical_activity',
 'num_BMI', 'num_HbA1c', 'num_weight',
 'num_insulin_sensitivity', 'num_sleep_hours', 'num_creatinine',
 'nom_gender_female', 'nom_gender_male',
 'nom_previous_medications_both',
 'nom_previous_medications_insulin',
 'nom_previous_medications_none', 'nom_previous_medications_oral',
 'ord_family_history', 'ord_food_intake'], dtype=object)`

In [48]: `df.columns`

Out[48]: `Index(['gender', 'age', 'family_history', 'glucose_level', 'physical_activity',
 'food_intake', 'previous_medications', 'BMI', 'HbA1c', 'weight',
 'insulin_sensitivity', 'sleep_hours', 'creatinine', 'Insulin'],
 dtype='object')`

In [49]: `gender = 'male'`
`age = 35`
`family_history = 'no'`
`glucose_level = 140.00`
`physical_activity = 2.00`
`food_intake = 'low'`
`previous_medications = 'none'`
`BMI = 30.00`
`HbA1c = 8.00`
`weight = 80.00`

```
insulin_sensitivity = 1.50
sleep_hours = 8.00
creatinine = 1.30

inputs = [[gender, age, family_history, glucose_level, physical_activity, food_intake, previous_medi
```

```
In [50]: input_df = pd.DataFrame(inputs, columns=df.columns[:-1])
# input_df = pd.DataFrame(inputs, columns=transformed_columns)
input_df
```

```
Out[50]:    gender  age  family_history  glucose_level  physical_activity  food_intake  previous_medi
0      male   35           no        140.0            2.0          low
```



```
In [51]: transformed_data = preprocessor.transform(input_df)
transformed_data
```

```
Out[51]: array([[-0.74454064,  0.10325266, -1.03306914,  0.14213074, -0.01675114,
       0.01360689,  1.73228436,  0.58525623,  1.03409572,  0.          ,
       1.          ,  0.          ,  0.          ,  1.          ,  0.          ,
       0.          ,  1.          ]])
```

```
In [53]: prediction = model.predict(transformed_data)[0]
prediction
```

```
Out[53]: 'steady'
```